

CANDIDATE_INSTRUCTIONS

AI Application Engineer - Take-Home Project

Overview

Welcome to the technical assessment! This project evaluates your ability to build AI-powered applications using modern development tools and AI coding assistants. You'll create a sales forecasting system with backend ML pipeline, REST API, and a functional web interface.

Time Allocation: 2 days

Tools Allowed: Claude Code, Cursor, or any modern development tools

Philosophy: Focus on functionality and clear explanations

Dataset

You've been provided with `ecommerce_sales_data.csv` containing 669 days of e-commerce sales data with the following columns:

- `date`: Daily timestamp (2023-01-01 to 2024-10-30)
- `daily_sales`: Daily sales revenue in USD
- `product_category`: Product category (Electronics)
- `marketing_spend`: Daily marketing budget in USD
- `day_of_week`: Day name (Monday-Sunday)
- `is_holiday`: Binary flag (1=holiday, 0=regular day)

Your Task

Build a sales forecasting system with the following components:

Part 1: Backend & ML Pipeline (100 points)

1.1 Data Processing Pipeline (15 points)

Create a data processing system:

- Data validation and cleaning
- Feature engineering (at least: lag features, rolling averages, date features)
- Basic data quality checks

Deliverables:

- Data processing functions (can be in single file or module)
- Brief documentation of features created

1.2 Machine Learning Model(s) (35 points)

Implement at least ONE working forecasting model:

- **Choose one approach:** Prophet, SARIMA, XGBoost, LSTM, or similar
- Proper train/test split with temporal ordering (no shuffling!)
- Model evaluation with appropriate metrics (MAPE, MAE, RMSE)
- Save trained model for reuse

Bonus (+10 points):

- Compare 2 different model types
- Simple hyperparameter tuning

Deliverables:

- Model training script or notebook
- Saved model file(s)
- Model evaluation results showing performance

1.3 REST API (30 points)

Build a functional API using FastAPI or Flask:

Core Endpoints (Required):

POST `/api/forecast`

- Input: {horizon: number of days}
- Output: {predictions: [{date, value}]}

GET `/api/historical`

- Output: historical sales data

GET `/api/metrics`

- Output: model performance metrics (MAPE, MAE, RMSE)

Nice to Have (Optional):

GET `/api/feature-importance`

- Output: **feature** importance **if** using tree-based model

POST `/api/upload`

- Upload **new** CSV data

Requirements:

- Basic input validation
- Error handling with appropriate HTTP status codes
- CORS enabled for frontend
- API documentation (Swagger auto-generated is fine)

Deliverables:

- Working API with at least 3 core endpoints
- README with API usage examples

1.4 Data Storage (10 points)

Simple data persistence:

- **SQLite is fine** (no need for PostgreSQL)
- Store: historical_sales, predictions, model_metrics
- Basic schema (can be simple tables)

Deliverables:

- Database setup script or initial schema
- Functions to read/write data

1.5 Testing & Code Quality (10 points)

Basic quality assurance:

- At least a few unit tests for critical functions
- API endpoint tests (even just 2-3 key ones)
- Clean, readable code with comments
- Basic error handling

Deliverables:

- Test file(s) with some coverage
- Code that runs without errors

Part 2: Frontend Web Application (60 points)

2.1 Functional Web Interface (40 points)

Build a simple web interface (React, Vue, Svelte, or even plain HTML+JS):

Required Views (Choose approach that works for you):

Option A: Single-Page Dashboard (Simpler):

- Display key metrics (total sales, average, model MAPE)
- Historical sales chart (line chart)
- Forecast chart showing predictions
- Simple form to generate new forecasts (input: number of days)
- Display historical data in table

Option B: Multi-Page App (More structure):

1. **Overview:** KPIs + historical chart + latest forecast
2. **Forecast:** Form to generate predictions + visualization
3. **Data View:** Table showing historical data + model metrics

Deliverables:

- Working web interface that connects to your API
- Charts display data correctly
- User can trigger forecast generation
- Basic responsiveness (doesn't break on mobile)

2.2 Data Visualization (15 points)

Functional charts (doesn't need to be fancy):

- Use any library: Chart.js, Plotly, Recharts, or even matplotlib-generated images
- Line chart for time series data
- Forecast should be visually distinct from historical data
- Tooltips showing values on hover (nice to have)
- Basic legend

Focus on:

- Charts render correctly
- Data is readable and understandable
- Clear labels and axes

2.3 Usability & Documentation (5 points)

Make it understandable:

- Clear labels and headings explaining what each section shows
- Loading indicator when fetching data
- Error messages if API calls fail

- Brief text explaining what the forecast means
- Instructions for using the interface

Deliverables:

- User can understand what they're looking at without guessing
- Basic error handling
- Simple, clean layout (CSS framework like Bootstrap/Tailwind is fine)

Part 3: Documentation & Deployment (20 points)

3.1 Documentation (15 points)

Clear project documentation:

README.md must include:

- Project overview (what it does)
- Technology stack used
- **Setup instructions** (step-by-step to run locally)
- How to run the backend API
- How to run the frontend
- API endpoint examples
- Model performance results (MAPE, MAE, RMSE)
- **Your approach:** Brief explanation of your ML model choice and why

Code Documentation:

- Comments explaining complex logic
- Brief docstrings for main functions
- Explanation of key design decisions

AI Usage Report (1 page):

- How you used AI coding assistants
- Examples where you modified AI suggestions
- What worked well / what didn't

Deliverables:

- Complete README that someone can follow to run your project
- Code is understandable with comments

3.2 Deployment (5 points)

Make it easy to run:

Minimum (choose one):

- Simple `docker-compose up` that starts everything
- OR clear instructions to run backend + frontend separately
- Environment variables documented

Nice to Have:

- Requirements.txt / package.json with all dependencies
- Shell script to set up environment
- Live demo URL (bonus +5 points)

Deliverables:

- Can run your project by following README
- Dependencies clearly listed

Bonus Features (Optional - 20+ extra points)

Only if you have time - these are nice to have:

Easy Wins:

- Multiple model comparison (show 2 models side by side) (+5)
- Feature importance visualization (+3)
- Data upload endpoint to add new data (+4)
- Confidence intervals on forecasts (+4)
- Download forecast as CSV (+2)

More Advanced:

- What-if analysis (adjust marketing spend, see impact) (+8)
- Anomaly detection highlighting (+6)
- Live demo deployed online (+5)
- Real-time model retraining trigger (+5)

Technical Requirements

Backend:

- **Language:** Python 3.8+
- **Framework:** FastAPI or Flask (your choice)
- **ML Libraries:** Choose what you're comfortable with

- Prophet, SARIMA (statsmodels), XGBoost, scikit-learn, etc.
- **Database:** SQLite is fine (simple is good)
- **Testing:** pytest with a few key tests

Frontend:

- **Framework:** React, Vue, Svelte, or plain HTML/JS - whatever you know
- **Styling:** Use any CSS framework (Bootstrap, Tailwind) or plain CSS
- **Charts:** Chart.js, Plotly, Recharts - pick one that's easy
- **HTTP Client:** Fetch API or Axios

Code Quality:

- Clean, readable code with comments
- Basic error handling
- No hardcoded credentials
- Clear variable names
- Some tests for critical paths

Suggested Project Structure

Keep it simple! Here's one way to organize:

```

sales-forecasting/
├── backend/
│   ├── api.py          # FastAPI/Flask app with endpoints
│   ├── data_processing.py # Feature engineering
│   ├── model.py         # ML model training and prediction
│   ├── database.py      # SQLite setup and queries
│   ├── requirements.txt # Python dependencies
│   └── tests/
│       └── test_api.py  # A few test files
├── frontend/
│   ├── index.html       # Main HTML file
│   ├── app.js           # JavaScript for API calls and UI
│   ├── style.css         # Styling
│   └── package.json      # If using npm packages
├── data/
│   └── ecommerce_sales_data.csv
├── models/              # Saved model files
│   └── forecast_model.pkl
└── README.md
└── docker-compose.yml    # Optional: if using Docker

```

Or organize however makes sense to you! This is just a suggestion.

Submission Guidelines

What to Submit:

1. Source Code:

- Complete backend and frontend codebases
- All configuration files
- Database migration scripts
- Docker configuration

2. Documentation:

- README.md with setup instructions
- ARCHITECTURE.md with design decisions
- API.md with endpoint documentation
- Code comments for complex logic

3. Tests:

- Backend unit and integration tests
- Test coverage report
- API testing collection (Postman/Insomnia)

4. Deployment:

- Docker setup with instructions
- OR live demo URL (highly recommended)
- Environment configuration examples

5. AI Usage Report (1-2 pages):

- How you used AI coding assistants
- Examples of prompts you used
- Instances where you modified/rejected AI suggestions
- What you learned from using AI tools

Submission Format:

Option A: GitHub Repository (Preferred)

- Create a private GitHub repository
- Add comprehensive README
- Share access with interviewer email

- Include git history (meaningful commits)

Option B: Zip File

- Compress entire project
- Include all files except node_modules, **pycache**, venv
- Email download link (Google Drive, Dropbox, etc.)

Before Submission Checklist:

- Code runs without errors locally
- All tests pass
- README has complete setup instructions
- Environment variables documented
- Docker setup works (or live demo accessible)
- API documentation is complete
- Frontend is responsive and polished
- No sensitive data or API keys in code
- Git history is clean (if using Git)
- AI usage documented

Evaluation Criteria

You will be assessed on a 180-point scale (200+ with bonuses):

Category	Points	Key Focus
Backend & ML	100	Working model, API, basic tests
Frontend	60	Functional interface, readable charts
Documentation	20	Clear README, explainable code
Bonus Features	20+	Extra polish, advanced features

What We're Looking For:

Priority 1 - Functionality (Must Have):

- ML model that works and beats naive baseline ($MAPE < 20\%$)
- API endpoints that return correct data
- Frontend that displays forecasts clearly
- Project runs by following your README

Priority 2 - Quality (Important):

- Clean, understandable code
- Basic error handling
- Some tests for main functions
- Good documentation of your approach

Priority 3 - Polish (Nice to Have):

- Beautiful UI
- Advanced features
- High test coverage
- Live deployment

Tips for Success

Time Management (2 days = ~16 hours):

- **Day 1 (8 hours):**
 - Hour 1-2: Data exploration, feature engineering
 - Hour 3-5: Train and evaluate ML model
 - Hour 6-8: Build basic API with 3 core endpoints
- **Day 2 (8 hours):**
 - Hour 1-3: Build frontend with basic chart
 - Hour 4-5: Connect frontend to API
 - Hour 6-7: Write tests, add error handling
 - Hour 8: Documentation, README, final testing

Strategy:

1. **Start Simple:** Get one model working before trying multiple
2. **MVP First:** End-to-end basic version, then add features
3. **Use AI Smart:** Generate boilerplate, but validate everything
4. **Test As You Go:** Don't wait until the end
5. **Document While Coding:** Comments and README as you build
6. **Prioritize:** Working > Perfect. Functionality > Beauty.

AI Tool Usage:

- Use Claude Code or Cursor to accelerate development
- Have AI generate boilerplate and scaffolding
- Validate all AI suggestions before accepting
- Use AI for code review and optimization

- Document interesting AI interactions
- Show critical thinking - don't blindly accept

Common Pitfalls to Avoid:

- **Overengineering:** Don't build a distributed system for 669 rows of data
- **Perfect is the enemy of good:** A working basic system beats a half-finished fancy one
- **Ignoring time:** Track your hours, don't spend all day tuning hyperparameters
- **No testing:** At least test your API endpoints manually
- **Poor README:** If we can't run it, we can't evaluate it
- **AI Blindness:** Don't copy-paste without understanding

Getting Started

1. **Review requirements thoroughly**
2. **Set up your development environment**
3. **Create project structure**
4. **Start with backend API + one model**
5. **Add frontend when backend is stable**
6. **Test and iterate**
7. **Deploy early, deploy often**
8. **Document throughout**

Resources

You may use:

- Official documentation for any libraries
- Stack Overflow and similar resources
- AI coding assistants (Claude Code, Cursor, GitHub Copilot)
- Online tutorials and guides

You may NOT:

- Copy complete solutions from others
- Use pre-built forecasting platforms
- Share assignment with others

Submission Deadline

Due: [Specific date/time will be provided]

Late submissions: Accepted with point deductions (10% per day)

We're excited to see what you build! This is your chance to showcase your full-stack AI engineering skills. Good luck!

Questions or need more time? Email us ASAP - we're reasonable about extensions for valid reasons.